

# RAG in Government: Deployments, Comparative Practice, and Societal Risk

## TL;DR

- Government RAG (retrieval-augmented generation) has moved from pilot to production at speed across the U.S. federal government (GSA's USAi, State's StateChat, GOV.UK Chat, France's Albert, Singapore's Pair, Estonia's Bürokratt), but the U.S. has simultaneously stripped some of the rights-impacting safeguards that other jurisdictions (EU AI Act, UK GDS) are doubling down on — creating a transatlantic governance divergence.
  - The clearest documented harms from public-sector LLM/automation systems to date come from *adjacent* systems (Australia's Robodebt, the Dutch toeslagenaffaire/SyRI, NYC's MyCity chatbot), each of which encoded existing administrative bias and weakened due process; these are direct cautionary blueprints for what poorly governed government RAG will produce.
  - The dominant risk pattern is not adversarial attack but **accountability laundering** — RAG outputs are treated by frontline staff and citizens as authoritative while liability is diffused across vendor, agency, and "the AI" — a pattern that the *Moffatt v. Air Canada* (2024) precedent flatly rejected in the private sector but that has no clean public-sector analogue in the United States.
- 

## Key Findings

1. **The U.S. federal RAG stack now has a single anchor.** USAi.Gov (launched August 14, 2025, by GSA) is the central evaluation and access platform for federal AI; combined with the OneGov procurement vehicle and DoD's four \$200M contracts to OpenAI, Anthropic, Google, and xAI, this concentrates federal RAG on a four-vendor model layer with Palantir, Microsoft, AWS, and Google dominating the integration layer.
2. **OMB policy has been substantially rewritten and partially softened.** The Biden-era M-24-10 and M-24-18 framework — which required AI impact assessments, bias mitigation, and notice-of-adverse-decisions for "rights-impacting AI" by December 1, 2024 — was replaced in April 2025 by M-25-21 and M-25-22, which collapse categories, drop election-infrastructure and likeness-replication from the high-impact list, and extend compliance deadlines to April 15, 2026.

3. **GAO is the most active institutional check.** GAO-25-107653 (July 29, 2025) documented federal AI use cases nearly doubling (571→1,110) and **generative AI use cases growing ninefold (32→282)** between 2023 and 2024, U.S. GAO with 35 open recommendations to 19 agencies.
  4. **The UK has the most rigorously evaluated public RAG system in the world.** GOV.UK Chat moved from 76% to 90% accuracy over two pilots involving 10,000+ users and 26,000 questions; all 508 jailbreak attempts were blocked; average latency 10.7s. Production rollout began in early 2026 on an Amazon Bedrock / Anthropic Claude stack.
  5. **Estonia, Singapore, and France illustrate three distinct successful models** — open-source interoperable network (Bürokratt), centralized productivity suite (Pair), and sovereign-stack workforce tool (Albert). The EU's posture is regulator-first (AI Office) rather than deployer-first.
  6. **Documented public-sector AI harms are concentrated in welfare/eligibility automation, not RAG specifically — yet.** SyRI/toeslagenaffaire (Netherlands, 35,000+ wrongly accused parents), Robodebt (Australia, 526,000 illegitimate debts, A\$1.8B settlement), and NYC MyCity (documented unlawful advice on housing, labor, and consumer law) are the empirical foundation for predicting RAG-specific harms.
  7. **Accountability law is sparse and asymmetric.** The Canadian *Moffatt v. Air Canada* (2024 BCCRT 149) ruling established that companies *are* liable for chatbot output — but no analogous doctrine constrains U.S. agencies, which can claim sovereign immunity and post disclaimers.
- 

## Details

### SECTION 1: RAG IN THE UNITED STATES GOVERNMENT

#### 1.1 The federal policy frame (2023–2026)

Federal RAG deployment sits inside a regulatory architecture that has been substantially rewritten twice in 24 months:

- **EO 14110 (Biden, Oct. 30, 2023)** established trustworthy-AI principles and directed OMB to issue binding agency guidance.
- **OMB M-24-10 (March 28, 2024)** required Chief AI Officers, AI use-case inventories, AI impact assessments, real-world performance testing, and — critically — bias and disparate-impact mitigation for “rights-impacting” AI, with a hard December 1, 2024 compliance-or-discontinue deadline.

- **OMB M-24-18 (Oct. 3, 2024)** added procurement-side requirements for “rights-impacting” and “safety-impacting” AI contracts.
- **EO 14179 (Trump, Jan. 23, 2025)** rescinded EO 14110 and reframed federal AI around “removing barriers” and “global AI dominance.”
- **OMB M-25-21 and M-25-22 (April 3, 2025)** replaced the Biden memos, collapsed “rights-impacting” and “safety-impacting” into a single “high-impact AI” category, omitted election infrastructure and likeness-replication from the high-impact list, and extended the compliance-or-shutdown deadline to April 15, 2026. Brookings analyst Brooke Tanner characterized the memos as “continuity” rather than dismantlement, but several civil-rights protections — including the express obligation to notify individuals of adverse AI-mediated decisions — were dropped.
- **EO 14319 (“Preventing Woke AI in the Federal Government,” July 23, 2025)** and a subsequent OMB memo from Director Russell Vought require vendors to attest that procured LLMs are “truth-seeking” and “ideologically neutral,” with eligibility and termination criteria tied to these attestations. A Stanford HAI policy brief, *“Toward Political Neutrality in AI”* (September 10, 2025), concluded that *“True political neutrality in AI is impossible, but there are practical approximations of neutrality that developers can implement across different levels of AI systems.”*

## 1.2 USAi.Gov (GSA, launched August 14, 2025)

USAi is the centerpiece of the current administration’s federal AI build-out. It is a FedRAMP-aligned, FISMA-Moderate, no-cost-to-agency evaluation suite providing:

- A **unified chatbot** front end across multiple commercial LLMs (initially OpenAI, Anthropic Claude, Google Gemini, Meta Llama; Grok added later; **Anthropic removed February 27, 2026** by directive of GSA Administrator Edward C. Forst after a White House order to “IMMEDIATELY CEASE all use of Anthropic’s technology.” [GSA](#) State Department’s StateChat similarly migrated off Claude by March 6, 2026, triggering Anthropic to file two lawsuits including one invoking the Federal Acquisition Supply Chain Security Act). [Defense One](#)
- **API access** for embedding models into agency pipelines.
- A **console** with usage analytics, model-comparison dashboards, and centralized logging — features GSA explicitly maps to M-25-21 and the June 2025 Executive Order on Trustworthy AI. [Nextgov.com](#)

By the end of 2025, **20+ agencies** had onboarded, with adoption data expected in early 2026 per GSA Chief AI Officer Zach Whitman. [FedTech Magazine](#) USAi was built on top of **GSai**, GSA’s internal chatbot launched [FedScoop](#) March 20, 2025; an internal GSA Office of

Evaluation Sciences evaluation found that of users who tried GSAi in its first five weeks, the median was only six prompts, 82% used only the default model, the top barriers were “inaccurate content” (33%) and “poor output quality” (33%), and 52% cited lack of workflow integration as the top drawback. [Gsa](#)

USAi sits alongside two parallel procurement moves that materially changed the federal RAG vendor landscape in 2025:

- DoD awarded **\$200 million ceiling contracts each to Anthropic, OpenAI, Google, and xAI** in July 2025. [Nextgov.com](#)
- Through GSA's **OneGov vehicle**, Anthropic and OpenAI offered their flagship models to federal agencies for **\$1 per agency** for a year; [FedScoop](#) xAI offered Grok for **\$0.42 per organization** through March 2027. [Globalgovernmentforum](#)

1.3 Named agency RAG deployments

| Agency | System                                                                                        | Function                                                                          | Corpus                            | Status                                                                  |
|--------|-----------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------|-----------------------------------|-------------------------------------------------------------------------|
| State  | <b>StateChat</b> (Palantir UI + Azure OpenAI; migrated to Claude, then off Claude March 2026) | Drafting, translation, summarization, performance-review support, panel selection | Internal cables, policy documents | Production; ~45,000–50,000 of 80,000 staff                              |
| State  | <b>FAM Search</b>                                                                             | RAG over the 25,000–page Foreign Affairs Manual                                   | FAM                               | Production (Oct. 2024)                                                  |
| State  | <b>Northstar</b>                                                                              | News ingestion/translation across 200+ countries, 100+ languages                  | ~1M news articles/day             | Production                                                              |
| State  | <b>StateInsight</b> (Palantir)                                                                | Structured + unstructured financial/performance Q&A                               | Internal financial systems        | Production                                                              |
| VA     | <b>VA GPT</b> and <b>Microsoft 365 Copilot Chat</b> in clinical settings                      | Clinical documentation, summarization                                             | EHR + clinical notes              | Production; <b>flagged by VA OIG Jan. 15, 2025</b> for lacking patient- |

|                           |                                                                                                                   |                                                    |                                                         |                                 |
|---------------------------|-------------------------------------------------------------------------------------------------------------------|----------------------------------------------------|---------------------------------------------------------|---------------------------------|
|                           | Military.com                                                                                                      |                                                    |                                                         | safety coordination             |
| <b>VA</b>                 | Chatbot/virtual agent on VA.gov                                                                                   | Citizen-facing benefits navigation                 | va.gov content                                          | Production                      |
| <b>DHS</b>                | 183 AI use cases (2024 GAO survey) including border-inspection support                                            | Mixed                                              | Production/pilot                                        |                                 |
| <b>HHS</b>                | 271 AI use cases (top in 2024 GAO survey); polio-outbreak detection RAG                                           | Scientific publications                            | Pilot                                                   |                                 |
| <b>NASA<br/>Goddard</b>   | Internal workforce chatbot (Hatamleh program)<br>Federal News Network                                             | Drafting, code, internal knowledge                 | Internal docs                                           | In development                  |
| <b>GSA</b>                | <b>GSAi</b> , PBS AI Chatbot, NCMMS AI Chatbot, Leasing Desk Guide Bot, OCFO Chatbot                              | Acquisition research, facilities, leasing guidance | FAR, Leasing Desk Guide, 20,000 daily NCMMS work orders | Mixed production/pre-deployment |
| <b>IRS /<br/>Treasury</b> | Early adopter per GSA accounts; specific RAG deployments not yet publicly inventoried at the same level of detail | –                                                  | –                                                       |                                 |
| <b>DoD</b>                | Multiple, exempted from public AI use-case inventory                                                              | –                                                  | –                                                       |                                 |

The **GAO's July 29, 2025 report (GAO-25-107653)** is the authoritative public count: across 11 inventoried agencies, total AI use cases grew from **571 in 2023 to 1,110 in 2024**, and **generative AI use cases increased ninefold from 32 to 282** in a single year. Gao Officials at 10 of 12 surveyed agencies cited "existing federal policy" as an adoption obstacle. DAU

## 1.4 State and local

- **New York City — MyCity Chatbot (Microsoft Azure-backed, October 2023 launch).** The signal failure case. Investigations by **The Markup** and **THE CITY (Colin Lecher, March 29 and April 2, 2024)** demonstrated the bot advising landlords they could discriminate against Section 8 voucher holders (illegal in NYC since 2008), telling employers they could pocket workers' tips (violates NY Labor Law §196-d), advising businesses they could refuse cash (illegal since 2020), [Envive](#) and suggesting that a restaurant could serve cheese nibbled on by a rat if customers were informed. [Museum of Failure](#) NYU's Julia Stoyanovich called the deployment **"reckless and irresponsible."** Mayor Eric Adams refused to take it down, calling public testing "how you iron out the kinks." [StateScoop](#) Build cost: roughly **\$600,000**. [Entrepreneur](#) Mayor Zohran Mamdani announced its termination on January 28, 2026, calling it "functionally unusable"; [THE CITY](#) the bot went offline February 5, 2026. [THE CITY](#) The OECD AI Incidents Observatory formally classified MyCity under "Misinformation" (Incident 714).
- **State-level deployments** are documented by NASCIO; State CIOs' program director **Amy Hille Glasscock** has framed RAG as the technique state agencies are adopting [StateTech Magazine](#) precisely to mitigate hallucination risk after MyCity. North Carolina (Governor Stein's Sept. 2, 2025 EO No. 24) [Coates' Canons](#) and other states have moved on AI governance frameworks.

## 1.5 Procurement patterns

Three clear patterns:

1. **A four-firm oligopoly at the model layer:** OpenAI, Anthropic (now politically constrained), Google, and xAI dominate, with Meta Llama as the open-weight option. Mistral has no comparable U.S. federal footprint, in contrast to its dominance in France.
2. **A consolidating systems-integration layer:** Palantir (State StateChat/StateInsight), Microsoft (NYC MyCity, broad Azure OpenAI distribution), AWS Bedrock, Google (USAi), Elastic, and a long tail of integrators (Riva Solutions, TechSur, Booz Allen).
3. **OneGov / USAi as a centralizing chokepoint:** GSA is converting individual agency procurements into a curated central catalog, with FedRAMP authorization explicitly being prioritized for AI vendors. This is the most consequential procurement shift since the cloud-first push of the 2010s — and is what made the Anthropic eviction logistically possible to execute in days.

## 1.6 Congressional oversight

The GAO has emerged as the most active institutional check:

- **GAO-25-107653 (July 29, 2025):** Generative AI use and management at federal

agencies.

- **GAO-25-107933 (Sept. 9, 2025):** Identified **94 government-wide AI requirements** in current law/EO/guidance and **10 executive-branch AI oversight groups.** U.S. GAO
- **GAO-25-107197 (May 19, 2025):** AI use and oversight in financial services regulation.
- **GAO-25-107172 (April 22, 2025):** Generative AI's environmental and human effects.
- **GAO-25-108412 (April 9, 2025):** Fraud and improper payments — AI data quality.
- **GAO-25-107435 (Dec. 18, 2024):** DHS risk-assessment guidance for AI in critical infrastructure.

Across these, GAO has issued **35 open recommendations to 19 agencies** as of mid-2025.

## 1.7 Documented federal failures

- **VA OIG Jan. 15, 2025 preliminary advisory memorandum:** VHA had authorized VA GPT and Microsoft 365 Copilot Chat for use with patient health information *without coordination with the National Center for Patient Safety*. The OIG explicitly identified “a potential patient safety risk” Military.com when AI omits or mischaracterizes information in clinical summarization.
- **Anthropic eviction and Anthropic v. United States (March 2026):** GSA's removal of Anthropic from USAi and the Multiple Award Schedule following a presidential directive GSA is itself a governance failure mode — vendors can be removed for political reasons in days, with no equivalent process protecting accuracy or civil-rights performance.
- **Federal employee experience data (GSA OES, 2025):** Inaccurate content and poor output quality were the top barriers to internal RAG adoption — a finding consistent with academic benchmarks of RAG on government corpora.

---

## SECTION 2: OTHER GOVERNMENTS USING RAG SUCCESSFULLY

### 2.1 United Kingdom — GOV.UK Chat (Government Digital Service)

The UK's deployment is the most rigorously evaluated public-sector RAG system in the world. Key facts:

- **Architecture:** RAG over ~100,000 GOV.UK pages (some sources say 80,000; GDS's December 2025 blog cites 100,000). Originally built on OpenAI GPT-4o/GPT-4o mini; Inside GOV.UK the production stack, per GDS's March 16, 2026 retrospective, now runs on **Amazon Bedrock with Anthropic Claude models.** UKAuthority

- **Evaluation scale:** Two scaled public pilots over 18 months; **10,000+ users, 26,000 questions** — “the government’s biggest public test of generative AI to date.”

UKAuthority

- **Measured outcomes (GDS retrospective, March 16, 2026):**

- User satisfaction: **64%** satisfied; **73%** of app users described it as useful.

UKAuthority

- Accuracy: improved from **76% to 90%** during the pilot. UKAuthority

- Average response latency: **10.7 seconds**. UKAuthority

- Of **508 jailbreak attempts**, all were blocked. UKAuthority

- In-scope answer rate raised to **88%** through clarifying follow-up questions.

UKAuthority

- **Rollout plan:** Production deployment via the GOV.UK app starting early 2026; site-wide rollout to follow. Inside GOV.UK GDS is openly exploring **agentic** evolution — modeled on Ukraine’s **Diia.ai**, which already issues income certificates on request. The Register

- **Why it counts as “success”:** A specific, audited corpus; published evaluation data; an explicit refusal pattern for out-of-scope queries; “Check this answer” source links beneath every response; Inside GOV.UK and the entire pilot is documented on the public GDS blog.

## 2.2 European Union institutions

- **GPT@EC** (launched October 22, 2024 as a pilot by DG DIGIT, building on JRC’s earlier **GPT@JRC**): internal corporate generative-AI tool for Commission staff, alternative to third-party tools, processes internal documents without external sharing.

European Commission Aligned with the **AI@EC Communication** (January 2024).

- **AI Office** within DG CONNECT, established January 24, 2024, supervises GPAI obligations under the **EU AI Act**, Wikipedia with ~125+ staff European Commission and the power to fine up to **3% of worldwide annual turnover or €15 million** Wikipedia for GPAI infringements.

- **GenAI4EU:** per the European Commission’s own GenAI4EU page, “*The EU is deploying close to €700 million specifically for generative AI through its GenAI4EU initiative,*” as part of the **AI Continent Action Plan** (April 9, 2025).

The EU’s posture is distinctive: **the AI Office is regulator, not deployer**. Where the U.S. has built a central LLM platform (USAi), the EU has built a central regulator and is letting member states deploy their own RAG systems under common rules.



## 2.3 Estonia — Bürokratt

- **Architecture:** Open-source, interoperable network of chatbots across public-sector institutions, [Interoperable Europe](#) available via website and (in development) voice. Originally built on Rasa [Kratid](#) NLP/BookStack; the system added **LLM + RAG** for accuracy-grounded answers in 2025, per RIA.
- **Adopters:** Estonian Tax and Customs Board (EMTA), Police and Border Guard, [RIA](#) Estonian Health Insurance Fund, Consumer Protection and Technical Regulatory Authority [RIA](#) (TTJA), National Library, Statistics Estonia, Defence Resources Agency, Rae Parish, Ministry of Justice's legal chatbot juristaitab.ee, [RIA](#) and others.
- **Cost:** Per CDO Ott Velsberg quoted by e-Estonia, *"The government has spent about €1.5 million on Bürokratt to date, and over the next four years, the total budget will be close to €13 million."* [e-Estonia](#) Marginal hosting cost per institution ~**€150/month plus LLM usage**. [Kratid](#)
- **Cross-border:** Designed for interoperability with **Finland's AuroraAI** [EGovReview](#) and discussions opened with **Belgium, Luxembourg**, and the **UK**.
- **Why this is the European model case:** UNESCO IRCAI has listed Bürokratt in its global Top 100 AI projects in multiple consecutive years. [Invest in Estonia](#) Estonian CDO Ott Velsberg has framed the system around explicit transparency, data tracking, and citizen consent. Earlier versions deliberately used NLP rather than LLMs because, per BBC reporting, "they are unlikely to give wrong or misleading answers" [Biometric Update](#) — a conservative posture only now being relaxed with controlled RAG.

## 2.4 France — Albert

- **Architecture:** Sovereign LLM stack built on Meta Llama 3.1 and Mistral models, exposed via Albert API (interministerial inference platform). [Dualmedia](#) Heavy use of RAG with knowledge-graph experimentation underway.
- **Deployments:** France Travail (the public job-seekers agency) — **ChatFT** (writing assistance) and **MatchFT** (adviser support) [Euronews](#) — already used by 10,000 agents. Ministry of Armed Forces partnership with Mistral AI via AMIAD. [Mistral AI](#) Tax agents using Albert for query handling, with human validation. [Globalgovernmentforum](#)
- **Scale ambition:** Minister Laurent Marcangeli announced a partnership with Mistral to "equip 10,000 civil servants" with **Le Chat** initially, ultimately targeting all **5.7 million French public servants**.
- **AI minister:** France appointed **Clara Chappaz** [Globalgovernmentforum](#) as Secretary of State for AI and Digitalisation following the 2024 elections.

## 2.5 Germany

- **Bundesdruckerei's Assistent.iQ:** in-house sovereign AI assistant for federal administration, runs on-premises or in a German-controlled cloud at BSI "high" baseline protection; designed for GDPR and EU AI Act compliance with sentence- and paragraph-level source citation. [Bundesdruckerei-gmbh](#)
- **PLAIN platform:** BMI-led data/AI infrastructure for federal administration, IaaS/PaaS/SaaS, with a "Community-Mandant" tenant introduced in 2024. [IT-Planungsrat](#)
- **BSI Kriterienkatalog (June 24, 2025):** BSI's binding criteria catalog for the integration of externally provided generative AI in federal applications. [BSI](#)
- **Bundestag Wissenschaftlicher Dienst (WD 7 - 3000 - 004/25):** published a structured analysis of the federal admin's AI use risks.

## 2.6 Netherlands, Nordics

- The Netherlands' use is shadowed by the **toeslagenaffaire** (see Section 3) and is now constrained by the **AI Verordening**. Active municipal AI experiments continue but under tight oversight from the **Autoriteit Persoonsgegevens**.
- Denmark's Agency for Digital Government (Digitaliseringsstyrelsen) and Sweden's DIGG are running pilots; **Finland's AuroraAI** is the closest analogue to Bürokratt and is being made interoperable with the Estonian system.

## 2.7 Singapore — Pair (GovTech / Open Government Products)

- **Suite:** Pair Chat, Pair Assistant (custom GPT-style builds), Pair Search (RAG over parliamentary debates, Supreme Court judgments, legislation), Pair Noms (meeting minutes). [United Nations Development Programme](#)
- **Adoption:** As of June 2026 reporting, more than half of Singapore's **150,000 public officers** use Pair; **60,000+ registered users, 20,000+ weekly active users, 10M+ messages exchanged**, [United Nations Development Programme](#) with UNDP's Singapore Policy Centre stating that the platform is *"saving officers an estimated 46% time spent on administrative tasks and raising public sector productivity."*  
[United Nations Development Programme](#)
- **Agency cases:** Singapore General Hospital (surgical antibiotic prophylaxis), Building and Construction Authority (Integrated Digital Delivery review), National Heart Centre Singapore (clinical text analysis). [Pair](#)
- **Governance:** Built with internet-separated architecture, no LLM provider logging,

Singapore is rolling out an **AI Agent Registry** (June 2026) ahead of broader agentic use.

## 2.8 Canada and Australia

- **Canada:** Treasury Board's Directive on Automated Decision-Making (2019, periodically updated) requires Algorithmic Impact Assessments for ADM at four risk levels. Generative AI guidance from TBS in 2023–2024. RAG deployments tend to be agency-internal (Statistics Canada, IRCC), and Canada's *Moffatt v. Air Canada* (2024 BCCRT 149) is the leading common-law precedent on chatbot liability (Section 3).
- **Australia:** Robodebt's shadow (Section 3) has slowed citizen-facing automation. The DTA's interim guidance on generative AI (2023), the National AI Capability Plan, and the Royal Commission report shape deployments. The DTA's whole-of-government **GovAI Chat** trial used Microsoft Copilot for public servants.

## 2.9 What "success" looks like — and the metrics actually used

- **Process metrics that have appeared in primary sources:** GOV.UK Chat (accuracy 76%→90%, 88% in-scope answer rate, 10.7s latency); Pair (46% admin-task time reduction per UNDP); USAi (no quantified outcome data yet, ramp-up only).
- **Governance markers:** published evaluations (UK), open source (Estonia), AI Act compliance / regulator capacity (EU), bilingual/source-cited outputs (Germany Assistent.iQ).
- **Public reception:** Estonia and Singapore enjoy high public trust; the UK is cautiously positive after extensive piloting; the U.S. case is more contested, with the MyCity failure casting a long shadow.

---

## SECTION 3: BROADER SOCIETAL HARM FROM GOVERNMENT RAG SYSTEMS

### 3.1 The pre-RAG cautionary corpus

The harms most likely to recur in RAG deployments are already documented in adjacent automated-decision systems:

- **The Netherlands — SyRI and the toeslagenaffaire.** The District Court of The Hague ruled SyRI (System Risico Indicatie) violated **Article 8 ECHR** on February 5, 2020, Taxadmin AI citing opacity, risk of discrimination against low-income neighborhoods, and inadequate data minimization. The court explicitly noted "the deployment of SyRI in 'problem areas' will discriminate and stigmatize the citizens." LOC The

**toeslagenaffaire** (childcare-allowance scandal) saw **35,000+ parents wrongly accused of fraud**, [Sage Journals](#) with nationality used as a risk indicator; the Rutte cabinet resigned January 15, 2021. [Taxadmin AI](#) Investigations by **Lighthouse Reports and Argos** later showed that a slightly modified SyRI continued in use in some Dutch cities [Lighthouse Reports](#) — illustrating the “lazy fix” problem.

- **Australia — Robodebt.** Operated 2016–2019; 526,000 debt notices to roughly **1 in every 50 Australians**; [Oxford Academic](#) **663 debt notices were issued to vulnerable people who died soon after**; [Wikipedia](#) Federal Court settlement of **A\$1.8 billion**; [Artificial Intelligence Incident Datab](#) **\$746M refunded, \$112M compensation.** [Oxford Academic](#) The Royal Commission (2023) made damning findings, with referrals for potential civil or criminal investigation. [Pursuit by the University of Melbourne](#) Net cost to taxpayers: **A\$565 million** for a scheme advertised as savings. [Pursuit by the University of Melbourne](#)
- **NYC MyCity** (Section 1.4): the closest direct analogue to RAG harm, because MyCity was a RAG-style chatbot over NYC business regulations.

These cases share a structural pattern that RAG will inherit: **administrative datasets encode the discriminatory practices of the agencies that produced them**, so a system that retrieves and synthesizes from them faithfully *will* reproduce that discrimination, often with more confidence and less appealability than a human caseworker.

### 3.2 Algorithmic bias as a feedback loop

RAG’s promise — grounding outputs in agency documents — is also its risk: **the corpus is the policy**, and historical policy documents encode historical injustice. The most explicit articulation of this is in the AI Now Institute’s “AI Now 2023 Landscape: Confronting Tech Power” (Amba Kak and Sarah Myers West, April 11, 2023), which warned: *“the dependence of government on Big Tech affordances came into particular focus during the height of the pandemic, when many national governments turned to Big Tech infrastructure, networks, and platforms for basic governance functions”* [AI Now Institute](#) — placing the burden of demonstrating non-harm on the public rather than vendors. AI Now’s 2025 follow-up, “Artificial Power,” argues this dependence has deepened.

Brookings’ work on **algorithmic exclusion** (Nicol Turner Lee and colleagues) makes the complementary point: **RAG systems can fail by omission as well as commission** — when the agency corpus simply lacks information about a population, the system silently returns nothing useful [Brookings](#) or fabricates. This is the specific failure mode for non-English speakers, people without recent benefit-claim history, and those whose situations aren’t covered by published guidance.

The AlgorithmWatch / AI Forensics joint investigation (final report December 15, 2023;

methodology spanned August–October 2023, 1,000+ queries (AI Forensics on Bavarian, Hessian, and Swiss elections) found that *“one third of Bing Chat’s answers to election-related questions contained factual errors. Errors include wrong election dates, outdated candidates, or even invented scandals concerning candidates.”* (AlgorithmWatch)

(AlgorithmWatch) AI Forensics’ Riccardo Angius added: *“It’s time we discredit referring to these mistakes as ‘hallucinations’. Our research exposes the much more intricate and structural occurrence of misleading factual errors in general-purpose LLMs and chatbots.”*

(AlgorithmWatch) The same architectural failure mode applies to RAG over government documents.

### 3.3 Automating bureaucratic denial and erosion of due process

The Special Rapporteur on extreme poverty and human rights, **Philip Alston**, framed this in his October 2019 report to the UN General Assembly (A/74/493) on the “digital welfare state”: *“There is a grave risk of stumbling zombie-like into a digital welfare dystopia.”*

(NYU School of Law) He documented that digitization “has very often been used to promote deep reductions in the overall welfare budget, a narrowing of the beneficiary pool, the elimination of some services, the introduction of demanding and intrusive forms of conditionality... and a complete reversal of the traditional notion that the state should be accountable to the individual.” (Diginomica) Human Rights Watch’s Amos Toh, summarizing the report, warned that “the automation of key welfare functions without sufficient transparency, due process, and accountability raises the specter of mass violations of welfare rights.” (Human Rights Watch) (AlgorithmWatch)

When RAG mediates eligibility — even in advisory form — three distinct due-process erosions follow:

1. **Hidden reasoning:** A frontline caseworker who relies on a RAG answer is effectively delegating reasoning to an opaque retrieval-and-synthesis process; the citizen cannot meaningfully contest a decision they cannot reconstruct.
2. **Disclaimer-laundering:** As MyCity demonstrated, a “beta” disclaimer paired with city-branded delivery produces the worst of both worlds — authority without accountability. The chatbot *itself* answered “Yes” when users asked if they could rely on it for professional business advice, contradicting its own warning banners. (Envive)
3. **Asymmetric appeals:** Agencies retain the right to revise the source corpus continuously; citizens get one frozen answer at the moment of interaction. *Moffatt v. Air Canada* (2024 BCCRT 149, Member Christopher C. Rivers) rejected this asymmetry in private law, holding: *“It should be obvious to Air Canada that it is responsible for all the information on its website. It makes no difference whether the information comes from a static page or a chatbot.”* (Springer) The tribunal also rejected the suggestion that the chatbot was “a separate legal entity that is responsible for its own actions,” (McCarthy)

calling that “a remarkable submission.” McCarthy Damages: C\$812.02. Envive No equivalent doctrine has been authoritatively established for government RAG in the U.S.; the *Westfall Act* and sovereign immunity make it materially harder to sue.

3.4 Language access, disability, and digital literacy

Government RAG systems are typically built on a corpus of formal administrative documents in the dominant language of the state. Three predictable equity failures follow:

- **Non-English speakers:** U.S. agencies are required by **EO 13166 (2000)** to provide meaningful access; a RAG system trained on English-only corpora that machine-translates outputs will produce subtle errors precisely where the law is most consequential (immigration, benefits, tax).
- **Disability access:** Voice/chat interfaces (Estonia’s Bürokratt is explicit about supporting hearing- and vision-impaired users) EGovReview can *expand* access if designed for it, but most U.S. RAG deployments have not yet been audited against Section 508 in the way agency websites have.
- **Low digital literacy:** GSA’s own internal-user evaluation showed that even highly literate federal employees mostly issued shallow prompts; citizens with less domain knowledge will fare worse, and the systems are not designed to detect when the user is misunderstanding the answer.

3.5 Accountability gaps

Who is liable when government RAG gives wrong guidance that harms a citizen? Four plausible defendants — and four plausible defenses:

| Defendant                                  | Plausible defense                                                                |
|--------------------------------------------|----------------------------------------------------------------------------------|
| Agency                                     | Sovereign immunity; disclaimers; “advice not decision” framing                   |
| Vendor (e.g., OpenAI, Microsoft, Palantir) | Indemnification clauses; “tools, not advice” framing; section 230-style theories |
| Integrator                                 | “We delivered what was specified”                                                |
| The “AI” itself                            | Not a legal person (the move Air Canada tried, and lost)                         |

ACLU Senior Policy Counsel Cody Venzke, responding to OMB’s revised guidance in 2025, observed: “*Before using AI to decide who gets a job, mortgage, or federal benefits — and more — federal agencies must first make sure that AI is up to the task, fair, and safe — and*

*discontinue it when it's not."* [American Civil Liberties Union](#) The ACLU specifically criticized the removal from M-25-21 of the *express obligation to notify individuals of adverse AI-mediated decisions*, calling the change a weakening of "transparency for individuals impacted by automated decisions." [American Civil Liberties Union](#)

EPIC has noted: *"Agencies have begun to publish guidance for AI in public benefits programs, but without more transparency and safeguards in place, these AI tools will continue to produce errors and biases that disproportionately harm marginalized communities."* [EPIC](#)

### 3.6 Chilling effects

There is suggestive but not yet rigorous evidence that citizen distrust of AI-mediated government services suppresses uptake of legitimate benefits — a "chilling effect" parallel to that documented around immigration data-sharing. Estonia has actively tried to inoculate against this by giving citizens a **data tracker tool** showing every entity that has accessed their data. [Government-transformation](#) The UK's GOV.UK Chat preserves explicit human-handoff paths and "check this answer at source" links beneath every response, [Inside GOV.UK](#) which research found users specifically appreciate. These are governance choices, not technical inevitabilities.

### 3.7 The civil-society research base

The most active civil-society research bodies on these themes:

- **AlgorithmWatch (Berlin/Europe):** Their joint investigation with **AI Forensics** (December 15, 2023, with follow-ups through 2024) on Microsoft's Bing Chat / Copilot — *"One third of Bing Chat's answers to election-related questions contained factual errors."* [AlgorithmWatch](#) [AlgorithmWatch](#) Senior researcher **Salvatore Romano** stated: *"Our research shows that malicious actors are not the only source of misinformation; general-purpose chatbots can be just as threatening to the information ecosystem."* [AlgorithmWatch](#) AlgorithmWatch's "Automating Society" report series tracks government ADM deployments across Europe, including German federal plans to use chatbots for citizen enquiries.
- **AI Now Institute (NYU):** "AI Now 2023 Landscape: Confronting Tech Power" (Kak and West, April 11, 2023) and "Artificial Power: 2025 Landscape Report" frame government RAG as part of a broader tech-power consolidation. AI Now Chief AI Scientist **Heidy Khlaaf** has been explicit that *"the existing guardrails for generative AI are deeply lacking, and it has been shown how easily compromised they are."* [AI Now Institute](#)
- **Brookings:** Brooke Tanner on OMB continuity; Nicol Turner Lee on algorithmic exclusion; the Center for Technology Innovation more broadly.

- **ACLU:** 2026 joint research report “Using AI to Make Sense of AI Policy” (with CNTR, January 8, 2026); Cody Venzke’s policy work on federal AI use.
- **EPIC:** ongoing work on government AI in public benefits.
- **Lighthouse Reports:** investigations on continued SyRI use in Dutch municipalities.
- **OECD AI Incidents Observatory:** structured incident database that has formally catalogued NYC MyCity (Incident 714, “Misinformation”), Robodebt (Incident 57), and others.

### 3.8 Documented measurable harms (consolidated)

| System                               | Jurisdiction        | Years     | Documented harm                                                                                                                                           |
|--------------------------------------|---------------------|-----------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|
| SyRI / Toeslagenaffaire              | Netherlands         | 2014–2021 | 35,000+ parents falsely flagged; loss of custody, bank closures, suicides; entire cabinet resigned                                                        |
| Robodebt                             | Australia           | 2016–2019 | 526,000 illegitimate debt notices; suicides; A\$1.8B settlement; A\$565M net loss to taxpayers                                                            |
| MyCity                               | New York City       | 2023–2026 | Documented unlawful advice on housing, labor, consumer law; ~\$600K build cost; OECD Incident 714                                                         |
| VA generative AI in clinical use     | U.S. federal        | 2024–     | Identified patient-safety risk by OIG; <a href="#">Military.com</a> no documented patient harm yet, but no formal coordination with patient-safety office |
| Bing Chat / Copilot election queries | Germany/Switzerland | 2023–2024 | One-third factual-error rate on elections <a href="#">AlgorithmWatch</a> per AlgorithmWatch/AI Forensics                                                  |

## Recommendations

### For federal CIOs / Chief AI Officers (next 6 months)

1. Treat USAi as an *evaluation* layer, not a *deployment* layer. Don’t put rights-impacting RAG into production solely on the basis of USAi pre-clearance.
2. Build agency-specific corpus audits *before* RAG deployment: who is over/under-



represented in the source documents, and what historical policy decisions are baked in?

3. Maintain an internal “high-impact AI” register that is **stricter** than M-25-21 — i.e., preserve the dropped categories (notice of adverse decisions, election-related uses) on a voluntary basis. Document the decision so successors can revisit.
4. Mandate “check this answer” source links on every public RAG output, on the GOV.UK Chat model.
5. Establish a 14-day takedown SLA for any RAG-mediated misinformation, with a public log.

### **For state and local governments**

1. Read the MyCity post-mortem carefully. Do not deploy public-facing RAG over regulatory content without (a) a domain-specific evaluation harness, (b) an independent audit by a body with subpoena power over the corpus, and (c) a kill switch and budget for taking the system down.
2. Adopt Canada’s Algorithmic Impact Assessment template as a floor, even where state law does not require it.

### **For procurement officers**

1. Require sentence-level provenance in any government RAG contract (the Bundesdruckerei Assistent.iQ standard).
2. Avoid single-vendor lock-in at the model layer; M-25-22’s “vendor-lock-in” language is a usable hook.
3. Build contractual obligations for the vendor to indemnify against discriminatory outputs documented in agency civil-rights audits.

### **For oversight bodies (Congress, GAO, IGs, courts)**

1. Establish a public, government-wide **AI incident registry** modeled on the OECD/AIIA database but with mandatory agency reporting on systems above a use-case threshold.
2. Codify a *Moffatt v. Air Canada* analogue: an explicit statutory duty that agencies are responsible for the substantive content of any AI system they deploy or brand.
3. Restore in statute the notice-of-adverse-decision obligation that was removed from M-25-21.

### **Benchmarks that would change the recommendations**

- If GAO’s 2026 report shows agencies hitting 90%+ on M-25-21 compliance by April 15, 2026, the most aggressive recommendations above can be relaxed.

- If documented incidents involving public-facing federal RAG remain below MyCity-scale through the end of 2026, the case for restraint weakens.
  - If any federal court adopts the Moffatt reasoning against a U.S. agency, the accountability gap closes materially.
- 

## Caveats

- USAi adoption data is preliminary (early 2026 expected); the GAO ninefold-growth figure is for *use cases inventoried*, not deployments at production scale.
- Several of the most striking historical harms (Robodebt, SyRI/toeslagenaffaire) are *not* RAG systems strictly — they are rule-based or classical ML risk-scoring systems. They are included because they are the most rigorously documented examples of the *governance failure modes* that RAG will inherit, and because the MyCity case demonstrates that the same patterns appear in RAG specifically.
- The political environment around U.S. federal AI is unstable; the Anthropic eviction and pending Anthropic v. United States litigation could meaningfully reshape the vendor landscape within months.
- “Success” metrics for European deployments are largely *self-reported by deploying agencies* (GDS, GovTech, OGP); independent audits are scarce. The UK’s GOV.UK Chat numbers are the most independently scrutinizable because of GDS’s open-blog evaluation methodology, but they are still GDS’s own numbers.
- This report focuses on text-RAG; multimodal and agentic extensions (Ukraine’s Diia.ai, Singapore’s emerging agent registry, the State Department’s roadmap) raise additional risks not covered here.
- Stanford HAI’s September 2025 brief is explicit that “true political neutrality in AI is impossible” — which means EO 14319 and Vought’s implementing memo set a standard that cannot be technically met, only attested to. This is a governance failure mode in itself.